

# OS Exercise 3: Building a Simple Filesystem

Operating Systems Reichman University

09.06.2025

## Submission Notes:

- **Submission Deadline:** Via Moodle by 07.07.2025. **No late submissions will be accepted!**
- **Collaboration:** This assignment is to be done in pairs (i.e., two students per team).
- **Submission Format:** Submit a single ZIP file containing all the requested files.
- **README:** Include a `README.txt` file at the root of the ZIP. This file must contain the submitting students' information. Each line should follow this format:  
  
    First name, Last name, Student ID
- **Working Environment:** This assignment requires the use of an Ubuntu OS with x86 64 bit architecture, be sure to work with the appropriate virtual machine/container/device.
- **Compilation Environment:** The checker uses `gcc13` on Ubuntu 24.04 LTS with the C17 standard (Default in GCC13).
- **Compilation and Execution:** Ensure your code compiles without errors and runs as expected.
  - The course staff will only support compilation via a docker container that uses the provided docker reference file.
  - You are provided with a reference compile script called "build.sh", you may modify it as needed but the tester is going to use it to compile your code, therefore all commands should be relative.
  - You are provided with an example main file (`main.c`) that you can use to a usage flow with regard to the FS that you are implementing, it is included in the reference build script.
  - You are provided with a reference "run.sh" script, that runs your compiled code, the tester is going to use it to run your code (with a different main).
  - You must ensure that after executing "build.sh" and then "run.sh" that the program runs correctly.
  - **You must submit build.sh and run.sh.**
- **Testing:** Always Write your own unit tests to validate your implementation.
  - you can also use the given test file: `testfilesystem.c` that contains a `main` function to test basic functionality. **Do not submit testfilesystem.c.**
  - You are given a reference script that compiles and runs `testfilesystem.c` called "compile\_and\_test.sh". **Do not submit compile\_and\_test.sh.**
  - If you do not pass all the provided tests your submission will get a failing grade.
- **Covered Material:** This assignment covers topics from File system Presentation Recitation.
- **Header File Modification - You can't modify the header file (including the structs)**

- **Dynamic Memory Allocation:** - Dynamic memory allocation (i.e malloc) is not allowed (and is not needed).
- **Simplifying Assumptions:** You may assume that only single threaded process will use your library, therefore, **thread-safety (mutual exclusion) considerations can be Omitted.**

## Introduction

In this project, you will implement a block-based filesystem in C, contained entirely within a single disk-image file. Your implementation, called "OnlyFiles", will be a simplified filesystem that supports file operations without directories. By building your own miniature filesystem—modeled conceptually on something like ext4—you'll gain hands-on experience with the core concepts of: low-level I/O, on-disk data structures, metadata management, and robust error handling.

The filesystem will store its data in a regular file (10MB in size), which acts as a "virtual disk". You will implement a C API that allows formatting the filesystem, mounting and unmounting it, creating and deleting files, and reading from and writing to files.

## Learning Objectives

By completing this assignment, you will:

- Understand how filesystems manage data at the block level
- Implement core filesystem operations (create, read, write, delete)
- Work with system-level I/O calls instead of buffered I/O
- Gain practical experience with filesystem metadata structures
- Practice error handling for filesystem operations

## Disk Layout

Your virtual disk (10MB) is divided into fixed-size blocks of 4KB each, resulting in a total of 2560 blocks. These blocks are organized into four main regions:

Region	Starting Block	Offset (bytes)	Size	Description
Superblock	0	0	4 KB (1 block)	Global filesystem metadata
Block Bitmap	1	4 KB	4 KB (1 block)	Tracks free/used blocks
Inode Table	2	8 KB	32 KB (8 blocks)	File metadata (256 inodes)
Data Blocks	10	40 KB	9.96 MB (2550 blocks)	File contents

Table 1: On-disk region layout (constants defined in `fs.h`).

## Filesystem Components

### Superblock

The superblock is the first block in the filesystem and contains critical metadata about the entire filesystem structure. It serves as the "table of contents" for your filesystem.

```

1 typedef struct {
2     int total_blocks;    // Total number of blocks (2560 for 10MB)
3     int block_size;      // Size of each block (4096 bytes)
4     int free_blocks;     // Number of available blocks
5     int total_inodes;    // Total number of inodes (256)
6     int free_inodes;     // Number of available inodes
7 } superblock;

```

Listing 1: Superblock Structure

When implementing the superblock:

- During `fs_format()`, initialize all fields to their appropriate values
- During `fs_mount()`, read the superblock to verify this is a valid filesystem
- Update `free_blocks` and `free_inodes` when allocating or freeing resources
- Always write the updated superblock back to disk after changing its values

## Block Bitmap

The block bitmap keeps track of which blocks are free (0) or in use (1). Each bit in the bitmap corresponds to one block in the filesystem.

```

1 // Bitmap is stored as: unsigned char bitmap[MAX_BLOCKS / 8];
2
3 // To mark block N as used
4 bitmap[N/8] |= (1 << (N%8));
5
6 // To mark block N as free
7 bitmap[N/8] &= ~(1 << (N%8));
8
9 // To check if block N is used
10 if (bitmap[N/8] & (1 << (N%8))) {
11     // Block is in use
12 }

```

Listing 2: Bitmap Operations

The bitmap helps you quickly:

- Find free blocks when allocating space for files
- Mark blocks as free when deleting files
- Determine if there's enough space available for a write operation

## Inode Table

The inode table contains metadata for each file in the filesystem. Each inode is 128 bytes in size, and your filesystem supports up to 256 inodes (files).

```

1 typedef struct {
2     int used;            // 1 if inode is active, 0 if free
3     char name[MAX_FILENAME]; // File name (up to 28 chars)
4     int size;            // File size in bytes
5     int blocks[MAX_DIRECT_BLOCKS]; // Block pointers (12 max)
6 } inode;

```

Listing 3: Inode Structure

Key operations with inodes:

- When creating a file, find a free inode and initialize it
- When writing to a file, update the size and block pointers
- When reading from a file, use block pointers to locate data
- When deleting a file, mark the inode as free and release its blocks

**Note:** With 12 direct block pointers and 4KB blocks, the maximum file size is 48KB.

## Data Blocks

The data blocks region stores the actual contents of files. Each file can use up to 12 blocks (as specified by `MAX_DIRECT_BLOCKS`), which means the maximum file size is 48KB.

## API Specification

### Filesystem Operations

`int fs_format(const char* disk_path)`

- **Description:** Creates and initializes a virtual disk file and prepares the filesystem.
- **Parameters:** `disk_path` - path where the virtual disk file will be created
- **Returns:**
  - 0 on success
  - -1 on failure (e.g., cannot create file)
- **Actions:**
  - If file exists, it will be overwritten
  - Creates a 10MB file
  - Initializes superblock, bitmap, and inode table
  - Sets all blocks as free except those used for metadata

`int fs_mount(const char* disk_path)`

- **Description:** Loads an existing filesystem from a virtual disk file.
- **Parameters:** `disk_path` - path to the virtual disk file
- **Returns:**
  - 0 on success
  - -1 on failure (e.g., file doesn't exist or invalid filesystem)
- **Actions:**
  - Opens the virtual disk file
  - Reads and validates the superblock
  - Loads necessary metadata into memory

`void fs_unmount()`

- **Description:** Ensures all pending changes are written to disk and closes the filesystem.
- **Returns:** None
- **Actions:**
  - Flushes any cached data to disk
  - Closes the virtual disk file

## File Operations

`int fs_create(const char* filename)`

- **Description:** Creates a new empty file in the filesystem.
- **Parameters:** `filename` - null-terminated string (max 28 chars excluding null)
- **Returns:**
  - 0 on success
  - -1 if file already exists
  - -2 if no free inodes available
  - -3 for other errors
- **Actions:**
  - Checks if filename already exists
  - Finds a free inode
  - Initializes the inode with the filename and zero size
  - Updates the superblock (decrease free\_inodes)

`int fs_delete(const char* filename)`

- **Description:** Removes a file and frees its blocks.
- **Parameters:** `filename` - null-terminated string
- **Returns:**
  - 0 on success
  - -1 if file doesn't exist
  - -2 for other errors
- **Actions:**
  - Finds the file's inode
  - Marks all of the file's blocks as free in the bitmap
  - Marks the inode as free
  - Updates the superblock (increase free\_blocks and free\_inodes)

`int fs_list(char filenames[][MAX_FILENAME], int max_files)`

- **Description:** Lists files in the filesystem.
- **Parameters:**
  - `filenames` - pre-allocated 2D array to receive filenames
  - `max_files` - maximum number of filenames to retrieve
- **Returns:** Number of files found (0 to `max_files`), or -1 on error
- **Actions:**
  - Scans the inode table for used inodes
  - Copies filenames to the provided array
  - Returns the number of files found

`int fs_write(const char* filename, const void* data, int size)`

- **Description:** Writes data to a file, overwriting any existing content.
- **Parameters:**
  - `filename` - target file
  - `data` - pointer to data buffer
  - `size` - number of bytes to write
- **Returns:**
  - 0 on success
  - -1 if file doesn't exist
  - -2 if out of space (not enough free blocks)
  - -3 for other errors
- **Actions:**
  - Finds the file's inode
  - Calculates how many blocks are needed
  - Frees any previously allocated blocks
  - Allocates new blocks as needed
  - Writes data to the allocated blocks
  - Updates the inode (size and block pointers)
  - Updates the bitmap and superblock

`int fs_read(const char* filename, void* buffer, int size)`

- **Description:** Reads file content into a buffer.
- **Parameters:**
  - `filename` - source file
  - `buffer` - pre-allocated buffer to receive data
  - `size` - buffer size in bytes
- **Returns:**

- Number of bytes read on success
- -1 if file doesn't exist
- -3 for other errors

- **Actions:**

- Finds the file's inode
- Determines how many bytes to read (minimum of file size and buffer size)
- Reads data from the file's blocks into the buffer
- Returns the number of bytes read

## Implementation Requirements

### System Call Usage

Your code must use low-level system calls to access the virtual disk:

```

1 // Opening the virtual disk
2 int disk_fd = open(disk_path, O_RDWR | O_CREAT, 0644);
3
4 // Reading a block
5 lseek(disk_fd, block_num * BLOCK_SIZE, SEEK_SET);
6 read(disk_fd, buffer, BLOCK_SIZE);
7
8 // Writing a block
9 lseek(disk_fd, block_num * BLOCK_SIZE, SEEK_SET);
10 write(disk_fd, buffer, BLOCK_SIZE);
11
12 // Closing the disk
13 close(disk_fd);

```

Listing 4: Required System Calls

Do **NOT** use:

- High-level stdio functions (`fopen`, `fread`, `fwrite`, etc.)
- Memory mapping (`mmap`)
- Dynamic memory allocation (`malloc`, `calloc`, etc.)

### Block Access Pattern

For all file operations, you should follow this general pattern:

1. Find the target file's inode (if operation requires an existing file)
2. Calculate which blocks are involved in the operation
3. Read or write those blocks using `lseek` to position and `read/write` for I/O
4. Update metadata (inodes, bitmap, superblock) as needed
5. Write changes back to disk

## Implementation Steps

We recommend implementing the functions in this order:

1. `fs_format` - Start by creating the disk and initializing structures
2. `fs_mount` and `fs_unmount` - Basic filesystem access
3. `fs_create` and `fs_list` - Simple file management
4. `fs_write` and `fs_read` - Data access
5. `fs_delete` - Resource cleanup

## Testing Your Implementation

The provided `main.c` includes a basic test that:

- Formats a virtual disk
- Mounts the filesystem
- Creates a file
- Writes data to the file
- Reads the data back
- Unmounts the filesystem

You should extend this with additional tests:

- Creating multiple files
- Writing files of different sizes
- Filling the filesystem to capacity
- Deleting files and reusing the space
- Testing error conditions (file not found, disk full)

## Sample Helper Functions

These are some useful helper functions to consider implementing:

```
1 // Find an inode by filename
2 int find_inode(const char* filename);
3
4 // Find a free inode
5 int find_free_inode();
6
7 // Find a free block
8 int find_free_block();
9
10 // Mark a block as used
11 void mark_block_used(int block_num);
12
13 // Mark a block as free
14 void mark_block_free(int block_num);
15
16 // Read an inode from disk
```



```

17 void read_inode(int inode_num, inode* target);
18
19 // Write an inode to disk
20 void write_inode(int inode_num, const inode* source);

```

Listing 5: Useful Helper Functions

## Common Pitfalls

- Forgetting to update the superblock after allocating/freeing resources
- Not writing changes back to disk (especially metadata)
- Buffer overflows with filenames
- Not handling partial reads/writes correctly
- Miscalculating block offsets
- Forgetting to update the bitmap when allocating/freeing blocks

## Final Notes

- **Read the Instructions Carefully:** Ensure that you adhere to all restrictions and requirements.
- **Testing:** Test your code thoroughly with various inputs to ensure correctness.
- **Code Quality:** Write clean, well-commented code to help the graders understand your implementation.
- **Error Handling:** Pay special attention to error cases and provide appropriate error codes.
- **Grading Reminder:** A happy grader is a merciful grader!

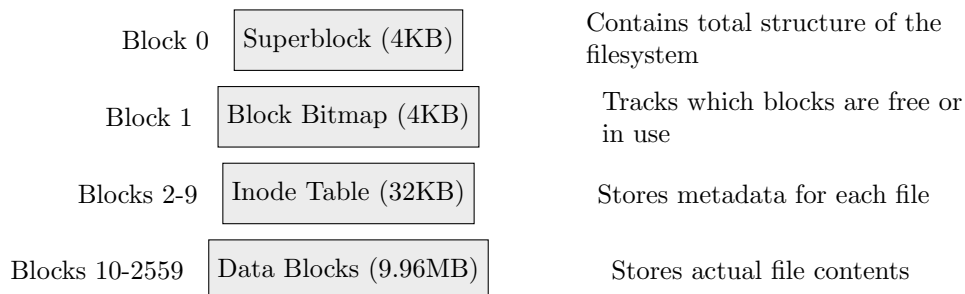


Figure 1: OnlyFiles Filesystem Structure

Good Luck!